

ExaNLA Survey Response Report

Generated on: 10/17/2025 at 12:48:23 PM

Submission Details

Library Name: FHI-aims
Version: 250822
Contact Name: Volker Blum
Email: volker.blum@duke.edu
Organization: Duke University

Selected NLA Operations

1. Quasi-Hermitian (BSE) Eigenvalue Problems
2. Symmetric/Hermitian Eigenvalue Problems
3. Matrix-Matrix Multiplication (GEMM)

1. Codes Information

Basic information about your application/simulation codes.

Library Name:
FHI-aims

Current Version:
250822

Contact Information:
Not specified

Name:
Volker Blum

Email:
volker.blum@duke.edu

Organization:
Duke University

Application Domain:
Not specified

What is the primary application domain of your codes?:
Materials Science

Materials Science:
Not specified

What are the main functionalities of your code?:
**Ground state DFT, Time-dependent DFT, Many-body perturbation theory (GW, BSE),
Molecular dynamics, Quantum transport, Excited-state dynamics, Crystal structure
prediction, Defect calculations, Surface science**

If you selected "Other", please specify::
Not specified

Climate/Weather Modeling:
Not specified

What are the main functionalities of your code?:
Not specified

If you selected "Other", please specify::
Not specified

Fluid Dynamics:
Not specified

What are the main functionalities of your code?:
Not specified

If you selected "Other", please specify::
Not specified

Other Domain Functions:
Not specified

What are the main functionalities of your code?:
Not specified

Use Case Information:
Not specified

Does your codes have multiple distinct use cases?:
Yes, multiple distinct use cases

Which use case are you describing in this submission?:
Groubd stte DFT

Library Description:
General purpose code for simulations in materials science and chemistry.

2. Matrix-Matrix Multiplication (GEMM)

Matrix-Matrix Multiplication (GEMM):
Yes

Matrix Properties:
Not specified

Matrix Structure:
Dense matrices, Distributed matrices

Matrix Distribution:
Block cyclic distribution (e.g., ScaLAPACK style)

Matrix Storage Format:
**Dense (column-major/row-major), Compressed Sparse Row (CSR/CRS),
Compressed Sparse Column (CSC/CCS), Multiple formats (conversion as needed)**

Which types of matrix multiplications do you perform?:
**Standard multiplication (A B), Transpose multiplication (A @ B
multiplication (A†B, AB†)**

Typical Dimensions:
Not specified

Matrix Size Range:
**Medium (100 - 1,000), Large (1,000 - 10,000), Very Large (10,000 - 100,000),
Extreme (> 100,000)**

Typical Matrix Shapes:
Square matrices (m "H n "H k)

Batch Size:
Not specified

Distributed-Memory NLA Library Usage:
Not specified

General Distributed Memory Libraries (CPU/GPU):
ScaLAPACK, ELPA, COSMA

Special/Advanced Implementations:
Not specified

Are there any NLA libraries you are interested in using (but have not yet adopted)?:
cuBLASMP (NVIDIA distributed-memory GPUs)

Future Requirements:
Not specified

Desired Features:
Improved sparse-dense multiplication, Hardware-specific optimizations, Other: simply, efficiency of all steps.

Benchmarking Requirements:
Not specified

Benchmark Input Types:
Real matrices from application workloads

Can You Provide Data or Mini-apps?:
Yes, matrices only

Scaling Requirements:
Strong scaling (fixed total problem size), Both strong and weak scaling needed

Working Precision:
Double precision (64-bit)

3. Standard Eigenvalue Problems ($Ax = \lambda x$)

Symmetric/Hermitian

Primary Use Cases:
Kohn–Sham equations (standard DFT), GW quasiparticle calculations, Bethe–Salpeter equation (Tamm-Dancoff approximation)

Matrix Properties and Structure:
Dense

Matrix Properties:
Not specified

Matrix Distribution:
Block cyclic distribution (e.g., ScaLAPACK style), Custom domain decomposition

Matrix Storage Format:
Dense (column-major/row-major), Compressed Sparse Row (CSR/CRS), Compressed Sparse Column (CSC/CCS)

Positive definiteness:
Varies depending on the problem

Eigenvalue distribution:
Varies

Problem Scale:
Medium (1,000 - 10,000), Large (10,000 - 100,000), Very Large (100,000 - 1,000,000), Extreme (> 1,000,000), Small (< 1,000)

Computation Requirements:
Not specified

Percentage of eigenvalues:
Varies

What to compute:
Varies

Eigenvalue location:
Varies

Required tolerance/precision:
Not specified

Residual tolerance type:
Absolute residual ($\|Ax - x\|$)

Absolute residual tolerance:
Machine precision

Relative residual tolerance:
Not specified

Hybrid residual tolerance:
Not specified

Orthogonality tolerance:
Machine precision

Working Precision:
Double precision (64-bit)

Workload Characteristics:
Not specified

Computation Pattern: capability or capacity:
Large-scale single problems (e.g., one large matrix at a time, using significant computational resources), Repeated similar-sized problems (e.g., time evolution or parameter sweeps)

Distributed-Memory NLA Library Usage:
Not specified

Distributed-Memory Dense Linear Algebra:
ScaLAPACK, ELPA

Iterative Eigensolvers:
Not specified

High-Level & Interface Libraries:
ELSI – Abstraction layer for eigenvalue solvers

Are there any NLA libraries you are interested in using (but have not yet adopted)?:
SLEPc, PRIMME, ChASE

Benchmarking Requirements:
Not specified

Benchmark Input Types:
Real matrices from application workloads

Can You Provide Data or Mini-apps?:
Yes, matrices only

Scaling Requirements:
Strong scaling (fixed total problem size)

Quasi-Hermitian (BSE) - $H\tilde{E} = E\tilde{E}$, where $H = \begin{pmatrix} A & B \\ -B^* & -A^* \end{pmatrix}$, $A =$

Matrix Properties and Structure:
Dense (standard full matrix)

Matrix Properties:
Not specified

Eigenvalue distribution:
Varies

Matrix scale/size:
Large (10,000 - 100,000)

Computation Requirements:

Not specified

Percentage of eigenvalues:

Varies

What to compute:

Varies

Eigenvalue location:

Varies

Required tolerance/precision:

Not specified

Residual tolerance type:

Absolute residual ($\|Ax - \lambda x\|$)

Absolute residual tolerance:

Machine precision

Relative residual tolerance:

Not specified

Hybrid residual tolerance:

Not specified

Orthogonality tolerance:

Machine precision

Working Precision:

Double precision (64-bit)

Workload Characteristics:

Not specified

Computation Pattern: capability or capacity:

Large-scale single problems (e.g., one large BSE matrix at a time, using significant computational resources)

Distributed-Memory NLA Library Usage:

Not specified

General Non-Hermitian Eigensolvers:

Solvers Targeting Full BSE Problems:

BSEPACK – Dedicated solver library for full BSE eigenproblems, ChASE – Chebyshev Accelerated Subspace Eigensolver, extended to BSE with customized filters and rayleigh-ritz

Are there any NLA libraries you are interested in using (but have not yet adopted)?:

ChASE

Benchmarking Requirements:

Not specified

Benchmark Input Types:

Real matrices from application workloads

Can You Provide Data or Mini-apps?:

Yes, matrices only

Scaling Requirements:

Strong scaling (fixed total problem size)

4. Generalized Eigenvalue Problems ($Ax = \lambda Bx$)

Symmetric/Hermitian A, SPD B

Matrix Structure:

A is dense, B is dense, A is banded, B is banded, Complex valued, Real valued

Reduction to Standard Eigenproblem (using B):

Not specified

Reduction to Standard Eigenproblem:

Yes, always

Reduction Method:

Cholesky factorization of B ($B = LL^T$ or $B = L^*L$)

Matrix Properties:

Not specified

Eigenvalue distribution:

Varies

Problem Scale:

Large (10,000 - 100,000)

Computation Requirements:

Not specified

Percentage of eigenvalues:

Varies

What to compute:

Varies

Eigenvalue location:

Varies

Required tolerance/precision:

Not specified

Residual tolerance type:

Absolute residual ($\|Ax - x\|$)

Absolute residual tolerance:

Machine precision

Relative residual tolerance:

Not specified

Hybrid residual tolerance:

Not specified

Orthogonality tolerance:

Machine precision

Working Precision:

Double precision (64-bit)

Workload Characteristics:

Not specified

Computation Pattern: capability or capacity:

Large-scale single problems (e.g., one large generalized eigenproblem at a time, using significant computational resources), Repeated similar-sized problems (e.g., time evolution or parameter sweeps)

Distributed-Memory NLA Library Usage:

Not specified

Distributed-Memory Dense Linear Algebra:

ELPA, ScaLAPACK

Iterative Eigensolvers:

Not specified

High-Level & Interface Libraries:

ELSI – Abstraction layer for eigenvalue solvers (e.g., used by SIESTA, FHI-aims)

Are there any NLA libraries you are interested in using (but have not yet adopted)?:

ChASE, DLA-Future

Benchmarking Requirements:

Not specified

Benchmark Input Types:

Real matrices from application workloads

Can You Provide Data or Mini-apps?:

Yes, matrices only

Scaling Requirements:

Strong scaling (fixed total problem size)